



PRICING AI FOR ENTERPRISE VALUE

A Commercial Architecture Framework for Scaling Purple Fabric

Customer Adaption x Value Realization x Intellect Economics

Sai Sreekar VVG

Enterprise AI Has Outgrown the Traditional Software Commercial Model

*Enterprise AI is moving from experimentation to enterprise-scale adoption.
“What will this cost to scale and what value will we get back?”*

MARKET SHIFT 1

AI Costs Are Variable

AI introduces operating costs behind each interaction, workflows and agent execution.

CHALLENGE

Usage can scale faster than budgets.

MARKET SHIFT 2

AI Value Is Not Automatic

Deployment does not automatically create business value.

CHALLENGE

Buyers increasingly demand measurable business impact.

MARKET SHIFT 3

AI Adoption Is Enterprise-Scale

AI is moving from isolated pilots to embedded workflows, platforms and autonomous agents.

CHALLENGE

A pricing model designed for experimentation may fail at scale.

The AI Commercial Dilemma: Maximize Adoption Without Sacrificing Economics

CUSTOMER FRICTION

When pricing is too rigid

- Unpredictable costs
- Fear of uncontrolled AI usage
- Difficulty forecasting future spend
- Uncertainty around ROI

CONSEQUENCE

Adoption slows. Pilots stall. Procurement pushes back. Deals take longer to close.

VALUE LEAKAGE

When pricing is too generic

- Value created beyond the price captured
- High-value automation treated like commodity software
- Compute and delivery risks absorbed by Purple Fabric
- Limited upside as customer value grows

CONSEQUENCE

Value is created but not captured. Margins compress. Commercial upside is left on the table.


The strategic challenge: reduce customer risk while allowing Purple Fabric to capture the value it creates.

Commercial Philosophy: Match the Commercial Model to the Value Created


Purple Fabric is not one product with one buying motion, it creates value across different layers of the AI stack.



THE USHO FRAMEWORK - Price the value pool not merely the underlying AI consumption.




One Purple Fabric Platform. Multiple Customer Value Pools.




Usage-Based Pricing: Pay for the AI Work Performed

BEST SUITED FOR



PURPLE FABRIC APPLICATION — CUSTOM TASK WORKLOADS



Usage converts variable AI consumption into a simple, measurable commercial unit.

Usage-Based Pricing Trades Commitment for Consumption

CUSTOMER FLEXIBILITY

Lower Upfront Commitment

Customer pays as the AI workload is executed.

- Lower barrier to pilot
- No need to predict long-term adoption
- Cost scales with actual activity

PURPLE FABRIC VOLUME EXPOSURE

Higher Volume Exposure

Revenue is directly linked to customer usage.

- Easier entry into uncertain use cases
- Revenue upside from adoption
- Direct linkage between activity and commercial value

Usage lowers the customer's psychological barrier to adoption , the commercial challenge is converting flexibility into sustainable economics.

A Usage-Based Proposal Has Four Commercial Layers

- 1

Annual License

Recurring commercial value for continued use of the Purple Fabric application and associated AI workload.
- 2

Implementation Fee

One-time enablement cost to move the solution from approved design to production readiness.
- 3

UAT / Bank-Managed Testing

Validation and acceptance cost.
- 4

Hypercare / Support

Post-go-live support.

The Annual License monetizes usage. The additional layers ensure complete delivery economics are covered.

The Annual License Calculator

WHAT THE ANNUAL LICENSE CALCULATOR CAPTURES

- AI work performed
- Workflow complexity
- Data and knowledge processed
- Infrastructure required
- Expected annual usage

COMMERCIAL UNIT

Workflow Execution

Technical Cost per Workflow × Expected Annual Workflow Volume

= *Annual Variable License Value*

TECHNICAL COST PER WORKFLOW

AI Model Consumption

Knowledge Consumption

Agent & Workflow Execution

API & Integrations

Data Processing

Infrastructure & Storage

Governance & Observability


The Customer sees only the Annual License Fees for using the AI application based on the usage like No. of Requests

The Commercial Manager sees only the 5 important variables to get an estimate cost for the proposal


AI Engineer / Technical Expert sees the complete backend for more specific use cases and technical calculation

Five Inputs Turn a Complex AI Workload Into a Budgetary Estimate

MANAGER INPUTS



OPERATING MODEL



Five simple commercial inputs. Detailed technical economics underneath.

Bioworld: From a Business Use Case to a Budgetary Commercial Estimate

THE MANAGER DASHBOARD

CUSTOMER

Bioworld Merchandising (India) Pvt. Ltd.

A major subsidiary of US-based Bioworld Merchandising, Inc., they are a global leader in brand growth and pop-culture licensing, operating primarily out of Gurugram and Noida. They hold licenses for over 100 global brands (including Disney, Marvel, and Warner Bros) and manage an extensive distribution network of over 2,100 points of sale in India.

USE CASE

The goal is to automate the verification process to ensure that all artwork listed in Centric PLM for a specific style correctly appears on the photographed merchandise by comparing it to reference images in the DMC Portal. The system will generate structured pass/fail reports with a Human-in-the-Loop (HITL) escalation path for low-confidence cases.

COMMERCIAL REQUIREMENT

50 requests / month — Tier A: 0–100 images · Tier B: 100–200 images — Avg. image size: 800 KB

Input	Tier A	Tier B	Technical Purpose
Deployment Model	PF Cloud / Client Cloud	PF Cloud / Client Cloud	Controls infrastructure economics
Task Complexity	3	3	PLM / DMC matching + HITL
Documents Processed	0	0	OCR used as unit proxy
OCR Volume	100	200	Represents image / vision volume
Storage Required	100 GB	200 GB	Image + audit capacity

The Underlying Workflow Economics behind the Calculator

Deployment	Tier A: 0–100 Images	Tier B: 100–200 Images
PF Cloud	₹175 / workflow	₹240 / workflow
Client Cloud	₹70 / workflow	₹140 / workflow

COST LOGIC BENEATH THE TABLE — PF CLOUD

PF Cloud — Tier A

LLM / Vision: ₹50

Knowledge: ₹5

Agent: ₹10

API: ₹5

Document Processing: ₹0

Storage: ₹0

Audit / Observability: ₹5

Infrastructure: ₹100

Total = ₹175 / workflow

PF Cloud — Tier B

LLM / Vision: ₹100

Knowledge: ₹10

Agent: ₹20

API: ₹10

Document Processing: ₹0

Storage: ₹0

Audit / Observability: ₹10

Infrastructure: ₹100

Total = ₹240 / workflow

Client Cloud

Variable AI / Knowledge / Agent / API costs remain, but PF Cloud infrastructure costs are removed.

Tier A = ₹70 / workflow

Tier B = ₹140 / workflow

The manager does not enter tokens, model calls, retrievals, compute, or API costs, those remain in the backend.

Bioworld: From Workflow Cost to Annual License Economics

ANNUAL VARIABLE COST - Cost per Workflow × Annual Volume Proposed Price to Bio World

Deployment	Tier A	Tier B
Cost / Workflow	₹175	₹240
Annual Volume (50 requests / month × 12 = 600 workflows / year)	600	600
Annual Total Cost - PF Cloud	₹2,25,000 / Year	₹2,70,000 / Year
Annual Total Cost - Client Cloud	₹42,000 / Year	₹84,000 / Year

One Time Setup Fee for Cloud

2.2. Deployment and Implementation Fee

Particulars	Fee in INR
PF Cloud tenant creation Fee	15,00,000
Implementation and Readiness to Go Live Fee	27,00,000

2.1. Annual License Fee for Bioworld

Particulars	Fee in INR
Annual License Fee (Up to 50 Request / month)	4,75,000 / Month
Additional Request	11,500 / Request
Additional Package of 50 Requests	4,75,000 / Month

Total Annual License Fee

Cost Component	Calculated Cost	Proposed Price	Variance / Premium
Annual Recurring Fee	₹2.70 Lakhs (Tier B)	₹57.00 Lakhs	+₹54.30 Lakhs (21.1x)
Cost per Workflow	₹240 / req	₹9,500 / req	+₹9,260 / req (39.6x)
Year 1 Total (Inc. Setup)	₹44.70 Lakhs	₹99.00 Lakhs	+₹54.30 Lakhs (+121%)

CALCULATED OVERHEAD ₹2.70 L/yr	PROPOSED PRICE ₹57.00 L/yr	GROSS PRICE MULTIPLE 21.1x Cost
-----------------------------------	-------------------------------	------------------------------------

Commercial Finding: Proposed recurring price (₹57 Lakhs/yr) represents a **21.1x markup (+2,011%)** over the maximum calculated execution cost (₹2.70 Lakhs/yr), yielding an estimated ~95.3% gross margin on compute operations.

Subscription-Based Pricing: Pay for Predictable Enterprise Access

CORE IDEA

The customer pays for predictable, recurring access to an enterprise AI capability, not for each individual AI execution.

Subscription monetizes ongoing enterprise capability rather than individual AI execution.

PRIMARY PURPLE FABRIC APPLICATION

Enterprise-Wide AI Assistants / Knowledge Access

Customer is buying:

- Ongoing capability
- Access
- Availability
- Predictable capacity

- User population is predictable
- AI interaction is recurring
- Enterprise knowledge is continuously accessed
- Customer wants predictable annual expenditure

CUSTOMER GETS

- Budget predictability
- No anxiety over individual usage
- Easier enterprise adoption
- Freedom to increase usage within agreed capacity

PURPLE FABRIC GETS

- Predictable recurring revenue
- Better recovery of platform costs
- Higher customer stickiness
- Ability to monetize enterprise-wide adoption

USAGE

“How much did I consume?”



SUBSCRIPTION

“What enterprise capability do I have access to?”

Subscription Uses the Same Commercial Foundation — With a Different License Logic

COMPONENTS

- 1. Annual License
- 2. Implementation Fee
- 3. UAT / Testing
- 4. Hypercare / Support
- 5. Other applicable commercial components

SUBSCRIPTION-BASED


The pricing is based on the Number of Seats or Users

The Customer sees only the Annual License Fees for using the AI application based on the seats like No.of Users

The Commercial Manager sees only the 5 important variables to get an estimate cost for the proposal

AI Engineer / Technical Expert sees the complete backend for more specific use cases and technical calculation

ANNUAL LICENSE CAPTURES




The Subscription Calculator Prices the Capacity Purple Fabric Must Provision

MANAGER-FACING INPUTS

- 1 User / Access Scale
- 2 AI Interaction Intensity
- 3 Knowledge / EKG Volume
- 4 Knowledge Change / Data Ingestion Rate
- 5 AI Application Complexity

BACKEND



Five business-level inputs drive a fully provisioned enterprise capacity plan.

Relationship Manager Assistant: From Enterprise Usage to Subscription Capacity

CUSTOMER REQUIREMENT

100 Relationship Managers

Each requires:

- Enterprise knowledge access
- Repeated daily AI interaction
- Shared organizational knowledge
- Predictable access
- Predictable annual expenditure

Input	Example
User / Access Scale	100 RMs
AI Interaction Intensity	High / recurring daily usage
Knowledge / EKG Volume	Enterprise knowledge base
Knowledge Change	Continuous / periodic updates
Application Complexity	Multi-agent RM Assistant

KEY POINT

The customer is not buying 100 × individual AI conversations — they are buying an enterprise AI capability available to 100 RMs within an agreed capacity envelope.

100 RMs > Access + Interaction + Knowledge + Ingestion + Orchestration capacity > Required Platform Capacity > Annual Subscription License

Relationship Manager Assistant: Backend Calculator Cost

Technical Cost Factor	PF Cloud Estimated Cost (₹/mo)	Client Cloud Estimated Cost (₹/mo)
1. User Access & Identity	₹15,000 – ₹30,000	₹0
2. Input / Prompt Processing	₹15,000 – ₹40,000	₹15,000 – ₹40,000
3. Knowledge Processing	₹10,000 – ₹30,000	₹10,000 – ₹30,000
4. EKG / Knowledge Storage	₹20,000 – ₹50,000	₹0
5. RAG / Vector Retrieval	₹25,000 – ₹60,000	₹25,000 – ₹60,000
6. LLM Inference	₹60,000 – ₹1,50,000	₹60,000 – ₹1,50,000
7. Agent / Orchestration	₹20,000 – ₹50,000	₹20,000 – ₹50,000
8. Tool & API Calling	₹15,000 – ₹50,000	₹0
9. File / Data Processing	₹5,000 – ₹20,000	₹5,000 – ₹20,000
10. Memory & Conversation	₹10,000 – ₹30,000	₹0
11. Compute & Infrastructure	₹50,000 – ₹1,50,000	₹0
12. Network / Data Transfer	₹5,000 – ₹20,000	₹0
13. Platform / Operational Overhead	₹30,000 – ₹80,000	₹0
TOTAL ESTIMATED MONTHLY COST	₹2,80,000 – ₹6,60,000 / month	₹1,35,000 – ₹3,50,000 / month

Hybrid Commercial Model: Price the Platform and the AI Work It Enables

Charge predictably for the platform. Charge proportionally for the AI workload it enables.

WHY ENTERPRISES NEED IT

- Give teams a persistent AI environment
- Build and orchestrate Digital Experts
- Deploy multiple AI applications
- Scale production workloads at different rates
- Maintain predictable platform economics while allowing variable execution

HYBRID COMBINES

Predictable Platform Capability

Variable AI Execution

PURPLE FABRIC

- Recurring base revenue
- Better cost recovery
- Upside from AI adoption
- Protection against uncontrolled consumption

CUSTOMER

- Predictable platform cost
- Defined capacity
- Flexibility to scale AI
- Clear commercial boundaries

Hybrid is not a compromise between Subscription and Usage — it is designed for customers buying the AI platform + the ability to scale what they build on it.

Purple Fabric Has Two Distinct Commercial Environments

ENVIRONMENT 1

Enterprise Chat / AI Access

Customer gets:

- Persistent enterprise AI environment
- AI interaction
- Organizational knowledge access
- Search / retrieval
- Documents and data
- AI assistants

PRIMARY ECONOMICS

Named users / workspace access · AI interaction capacity · Knowledge / EKG capacity · Storage · Platform capabilities

COMMERCIAL MODEL: SUBSCRIPTION

ENVIRONMENT 2

Digital Expert Build & Production

Customer gets:

- Digital Expert development
- Agent orchestration
- Tool / API connections
- Production workflows
- Business task execution at scale

PRIMARY ECONOMICS

Digital Experts · Workflow execution · Agent activity · API / tool consumption · Data processing · AI execution

COMMERCIAL MODEL: USAGE / CAPACITY

Hybrid Pricing Combines Predictable Platform Economics With Scalable AI Consumption

1. Platform / Subscription Capacity

Recurring base commercial.

- Named Designer Users
- Enterprise AI / Chat access
- Knowledge / EKG capacity
- Platform storage
- Agreed AI interaction capacity

PRINCIPLE

Pay for the platform capacity required to operate the enterprise AI environment.

2. Production Digital Expert Capacity

Variable / scale-linked commercial.

- Production Digital Experts
- Workflow execution capacity
- Agent / tool activity
- API consumption
- Data processing

PRINCIPLE

Pay according to the production AI capacity the enterprise requires.

3. Additional Capacity

When consumption exceeds baseline.

- AI capacity
- Knowledge / storage
- OCR / processing
- Infrastructure
- Execution capacity

PRINCIPLE

Base commitment + scale-linked expansion.

Platform Subscription

+

Production AI Capacity

+

Additional Consumption /
Overages

+

Implementation / UAT /
Hypercare

=

Total Commercial Value

Example: Enterprise AI Platform + Production Digital Experts

Sample Proposal - Metrics for the client

Platform Capacity

50 Designer Users

- Teams build and orchestrate AI applications
- Enterprise Chat / AI access
- Shared Knowledge / EKG
- Centralized governance

- Up to 50 named designer users. (Designer is a technical user who accesses PF platform to create agents and use cases).
- Up to 100 AI Agents.
- Up to 156 Million PF Tokens (defined herein below) per year. (Note: PF Tokens act as a unified AI credit covering all model usage - embedding, input processing, and output generation).
- Up to 200 GB of total knowledge storage, covering source documents, vectors, and indexes.
- OCR (Embedded 3rd party tool) usage: Up to 10,000 pages / month.
- Annual License Fee includes maintenance services L3 support from Intellect Manufacturing Centers.


Production AI Capacity (If the client wants to have both Developer and Production Setup Together)

Up to 100 Digital Experts

Multiple business workflows
API / tool integrations
Variable execution volumes
Production-scale AI workloads

OR

Developer Capacity (If the Client wants to opt for a seperate setup for Production and Development)



Outcome-Based Pricing: Price the Business Result

WHAT IS OUTCOME-BASED?

Customer pays based on a pre-agreed business outcome rather than underlying AI usage.

BEST SUITED FOR

Critical, repeatable workflows where the outcome can be:

- Clearly defined
- Objectively measured
- Attributed to Purple Fabric

WHEN TO USE

- Business-critical workflow
- Measurable outcome
- Meaningful delivery control
- Clear economic value
- Baseline exists

Claims Processing

Successful claim resolution

Customer Support

Resolving customer queries


Dimension	Customer	Purple Fabric
Risk	Lower upfront commercial risk	Higher delivery / performance risk
Psychology	<i>"I pay when I see the result."</i>	<i>"If we create greater value, we capture greater value."</i>
Strategic Intent	Align spend directly to business impact	Monetize the economic value created
Advantage	Higher confidence in ROI	Higher revenue potential and differentiation
Key Concern	Attribution and outcome measurement	Delivery control and margin protection

STRATEGIC DIFFERENCE
Usage VS Outcome

The Cost calculation remains the same as Usage Framework

Usage measures what Purple Fabric does.


Outcome measures what Purple Fabric achieves.



GUARDRAILS — outcome must be measurable, attributable, agreed upfront, within reasonable Purple Fabric influence, and protected by minimum commercial commitments where required.

Purple Fabric should not have one universal price model.

One platform. Multiple value pools. One disciplined commercial architecture.



Detailed Reports for Appendix